Project Presentation

Mi Yon Kim, Yufei Tao

Project Overview

- Dataset: Freeway data

- Tools & Languages:
  - MongoDB Compass
  - MongoDB Database
  - Python(Pymongo)

Data model

- Original model(RDBMS)

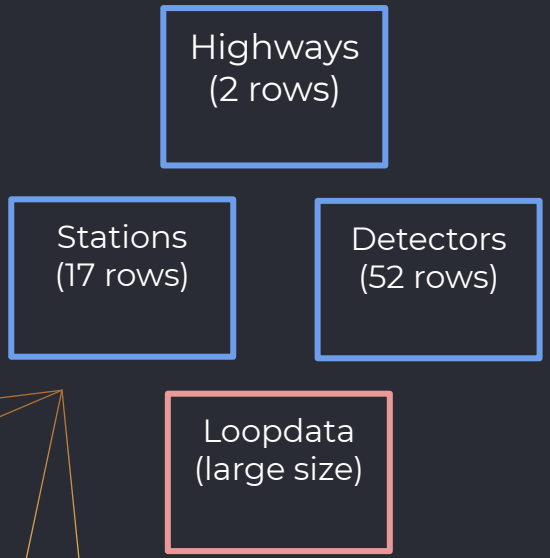

VS

- Our model(NoSQL)

  - **Detectors**
    - Combined 3 tables(Highway, Stations, Detectors)
    - Embedded documents of Highways and Stations.
    - Contains duplicated data however it can be minor because data is small

  - **Loopdata**
    - Added locationtext from the original model.

Data model

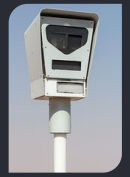


```
_id: ObjectId("5fced998e5d0fb0e54a4e125")
detectorid: 1345
milepost: 14.32
locationtext: "Sunnyside NB"
detectorclass: 1
lanenumber: 1
schema: "1.0"
station: Object
    stationid: 1045
    upstream: 0
    downstream: 1046
    stationclass: 1
    numberlanes: 4
    latlon: "45.43324,-122.565775"
    length: 0.94
highway: Object
    highwayid: 3
    shortdirection: "N"
    direction: "NORTH"
    highwayname: "I-205"
```


Detectors


```
_id: ObjectId("5fced9b3e5d0fb0e54a4e159")
detectorid: 1345
starttime: "2011-09-15 00:01:40-07"
volume: 1
speed: 47
occupancy: 0
status: 3
dqflags: 0
schema: "1.0"
locationtext: "Sunnyside NB"
```


Loopdata

ETL

- Implemented with python
- Drop loopdata which has the Null or zero value in the speed field
- Issues

  - Some data in Detector's locationtext and Station's locationtext didn't match
  (e.g. Columbia to I-205 NB(Station table) vs I-205 NB at Columbia(Detectors))

  - Some detectorid in Loopdata doesn't included in Detectors table.
  (e.g. Loopdata with 1350 detectorid vs 1350 doesn't show up in Detectors table)

  - Solution 1: Modified the csv file

  - Solution 2: Automate through script

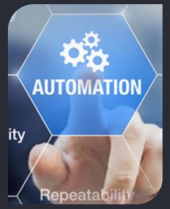

Query Examples

Q1: Count low speeds and high speeds:
Find the number of speeds < 5 mph or > 80 mph in the data set.

```
[ { 'Number of speeds < 5 mph and > 80 mph: ': 130182 } ]
```

```
db.Loopdata.aggregate(
  [
    {
      $match: {
        $or: [
          {
            speed: {
            $lt: 5
            }
          },
          {
            speed: {
            $gt: 80
            }
          }
        ]
      }
    },
    {
      $count: "Number of speeds < 5 mph and > 80 mph: "
    }
  ]
)
```

Query Examples

Q2: Volume: Find the total volume for the station Foster NB on Sept 15, 2011.

```
[ { _id: 0, 'Total volume:': 49870 } ]
```

```
# use $options:'i' to make the query case-insensitive
query = db.Loopdata.aggregate(
  [
    {
      "$match": {
        "$and": [
          {"starttime" : {"$regex": "2011-09-15.*"}
          },
          {"locationtext" : "Foster NB"
          }
        ]
      }
    },
    {
      "$group": {
        "_id": 0,
        "Total volume:": {
          "$sum": "$volume"
        }
      }
    }
  ]
)
print(list(query))
```

Query Examples

Q5: Route Finding: Find a route from Johnson Creek to Columbia Blvd on I-205 NB using the upstream and downstream fields.

```
# Get all of the NB data
results = list(col_detectors.find({"locationtext": {"$regex": direction+"$", "$options":'i'}}, {"station.stationid":1, "locationtext":

# save downstream of start point to find the next point
d_stream = start_doc["station"]["downstream"]
# Add start point
routes.append(start_doc["locationtext"])
for i in range(len(results)):
    depart = False
    for j in range(len(results)):
        # When finding previous point's downstream == stationid, add locationtext to routes
        # update the downstream
        # if updated downstream == end point's, ends of travel
        if d_stream == results[j]["station"]["stationid"]:
            routes.append(results[j]["locationtext"])
            d_stream = results[j]["station"]["downstream"]
            if d_stream == end_doc["station"]["stationid"]:
                depart = True
                break
            break
    if(depart == True):
        break

# Add end point
routes.append(end_doc["locationtext"])
```

```
(base) C:\gitRepo\CloudCluster\pdx-cs-cloud-cluster\queries>python3 query5.py
Johnson Cr NB
Foster NB
Powell to I-205 NB
Division NB
Glisan to I-205 NB
Columbia to I-205 NB
```

Query Examples

Q6: Update: Change the milepost of the Foster NB station to 22.6.

```
results = col_detectors.find({"locationtext": "Foster NB"})
milepost = "18.1"

print("\n\n---- Display the data before updating milepost ----")
for r in results:
    print("milepost:", r["milepost"])
    col_detectors.update_one({"milepost":r.get("milepost")}, {"$set":{"milepost": milepost}})

results = col_detectors.find({"locationtext": "Foster NB"})

print("\n\n---- Display the data after updating milepost to " + milepost + " --
for r in results:
    print("milepost", r["milepost"])
```

```
---- Display the data before updating milepost ----
milepost: 18.1
milepost: 18.1
milepost: 18.1


---- Display the data after updating milepost to 22.6 ----
milepost 22.6
milepost 22.6
milepost 22.6
```

Demo

Critique

- Changed Loopdata model
  - Reason: MongoDB document size limit: 16mb

```
loopdata {
    _id
    schema
    detectorid
    data: [
      {
        starttime
        volume
        speed
        occupancy
        status
        dqflags
      }
    ]

    locationtext
    num_lowspeed
    num_highspeed
}
```

Previous loopdata

Pros
- Improve performance on query #1
- No duplicate data

Cons
- Complicated
- Doesn't fit in MongoDB System

```
_id: ObjectId("5fcdd22384a63b4ed4be430d")
detectorid: "1345"
starttime: "2011-09-15 00:01:40-07"
volume: 1
speed: 47
occupancy: 0
status: 3
dqflags: 0
schema: "1.0"
locationtext: "Sunnyside NB"
```

current loopdata

Pros
- Simpler
- Simpler to make queries

Cons
- Duplicate data
- Bad performance on query #1 if the data size gets bigger

Critique

- Advice
  - Spend more time on researching new system before you design

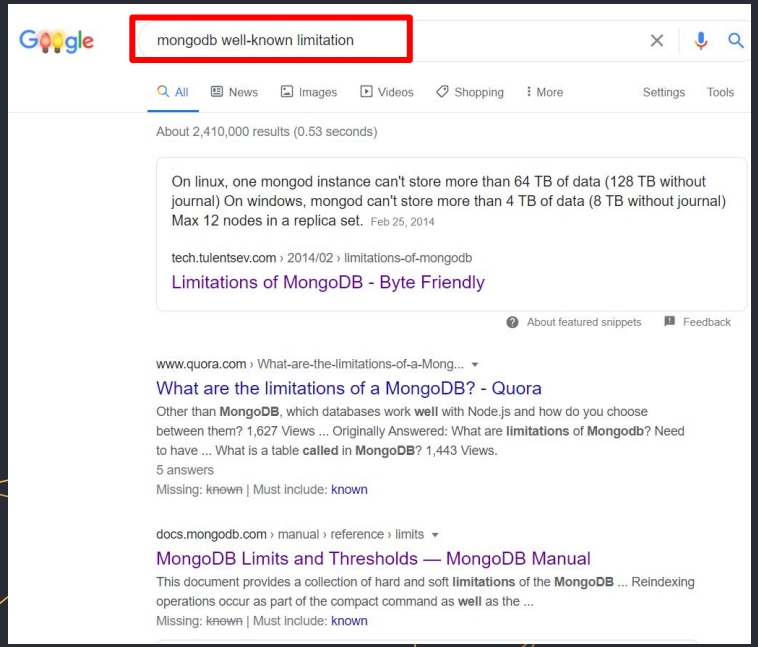

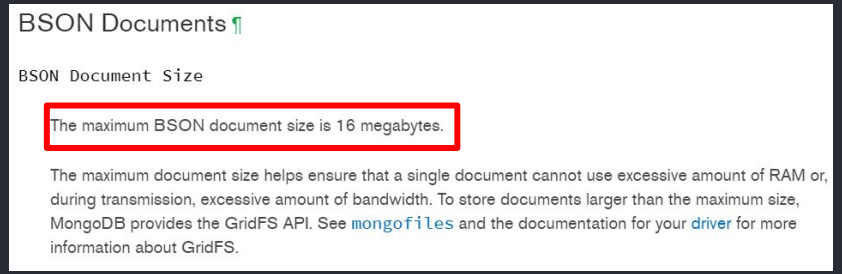

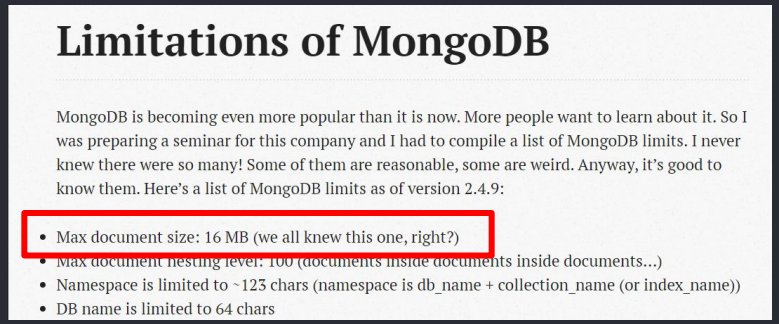

Lessons learned

1. Declare each attributes with the proper data type.

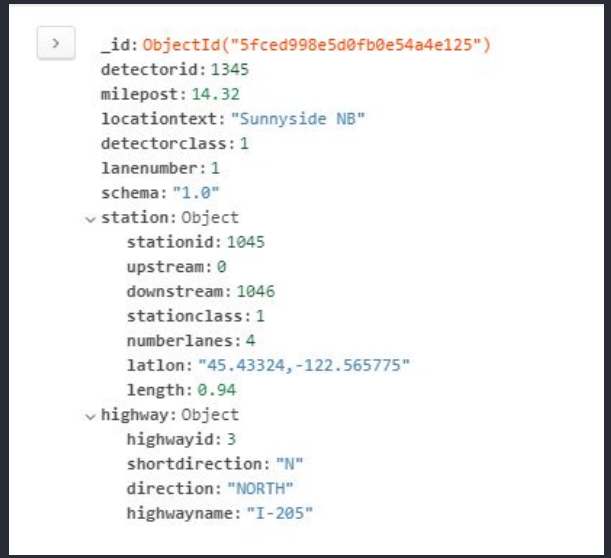

Lessons learned

2. Writing query purely using MongoDB could make things look messy. But they are easy to understand.

```
cursor = db.Loopdata.aggregate([{"$match": {"$or": [{"speed": {"$lt": 5}}, {"speed": {"$gt": 80}}]}},{"$count": "Number of speeds < 5 mph and > 80 mph: "}])
```

```
query = db.Loopdata.aggregate(
  [
    {
      "$match": {
        "$and": [
          {"starttime" : {"$regex": "2011-09-15.*"}
          },
          {"locationtext" : "Foster NB"
          }
        ]
      }
    },
    {
      "$group": {
        "_id": 0,
        "Total volume:": {
          "$sum": "$volume"
        }
      }
    }
  ]
)
```

VS

```
import pymongo
import os
from pymongo import MongoClient
from dotenv import load_dotenv
load_dotenv()
MONGO_HOST = os.getenv("MONGO_HOST")
MONGO_DB = os.getenv("MONGO_DB")

# connect to cluster
cluster = MongoClient(MONGO_HOST, 27017)
db = cluster[MONGO_DB]

# Get collection for query
col_detectors = db["Detectors"]

results = col_detectors.find({"locationtext": "Foster NB"})
milepost = "18.1"

print("\n\n---- Display the data before updating milepost ----")
for r in results:
    print("milepost:", r["milepost"])
    col_detectors.update_one({"milepost":r.get("milepost")}, {"$set":{"milepost": milepost}})

results = col_detectors.find({"locationtext": "Foster NB"})

print("\n\n---- Display the data after updating milepost to " + milepost + " ----")
for r in results:
    print("milepost", r["milepost"])
```

Lessons learned

3. Solve complicated questions with combination of python and queries.

```
...
import pymongo
import os
from pymongo import MongoClient
from dotenv import load_dotenv
load_dotenv()
MONGO_HOST = os.getenv("MONGO_HOST")
MONGO_DB = os.getenv("MONGO_DB")

# connect to cluster
cluster = MongoClient(MONGO_HOST, 27017)
db = cluster[MONGO_DB]

# Get collection for query
col_detectors = db["Detectors"]

results = col_detectors.find({"locationtext": "Foster NB"})
milepost = "18.1"

print("\n\n---- Display the data before updating milepost ----")
for r in results:
    print("milepost:", r["milepost"])
    col_detectors.update_one({"milepost":r.get("milepost")}, {"$set":{"milepost": milepost}})

results = col_detectors.find({"locationtext": "Foster NB"})

print("\n\n---- Display the data after updating milepost to " + milepost + " ----")
for r in results:
    print("milepost", r["milepost"])
```

Q & A

Appendix

Project github link: youn0125/pdx-cs-cloud-cluster (github.com)

Appendix

Q1: Count low speeds and high speeds:
Find the number of speeds < 5 mph and > 80 mph in the data set.

```
[ { 'Number of speeds < 5 mph and > 80 mph: ': 130182 } ]
```

```
db.Loopdata.aggregate(
  [
    {
      $match: {
        $or: [
          {
            speed: {
            $lt: 5
            }
          },
          {
            speed: {
            $gt: 80
            }
          }
        ]
      }
    },
    {
      $count: "Number of speeds < 5 mph and > 80 mph: "
    }
  ]
)
```

Appendix

Q2: Volume: Find the total volume for the station Foster NB for Sept 15, 2011.

```
[ { _id: 0, 'Total volume:': 49870 } ]
```

```
# use $options:'i' to make the query case-insensitive
query = db.Loopdata.aggregate(
  [
    {
      "$match": {
        "$and": [
          {"starttime" : {"$regex": "2011-09-15.*"}
          },
          {"locationtext" : "Foster NB"
          }
        ]
      }
    },
    {
      "$group": {
        "_id": 0,
        "Total volume:": {
          "$sum": "$volume"
        }
      }
    }
  ]
)
print(list(query))
```

Appendix

Q5. Route Finding: Find a route from Johnson Creek to Columbia Blvd on I-205 NB using the upstream and downstream fields.

```
# Get all of the NB data
results = list(col_detectors.find({"locationtext": {"$regex": direction+"$", "$options":'i'}}, {"station.stationid":1, "locationtext":

# save downstream of start point to find the next point
d_stream = start_doc["station"]["downstream"]
# Add start point
routes.append(start_doc["locationtext"])
for i in range(len(results)):
    depart = False
    for j in range(len(results)):
        # When finding previous point's downstream == stationid, add locationtext to routes
        # update the downstream
        # if updated downstream == end point's, ends of travel
        if d_stream == results[j]["station"]["stationid"]:
            routes.append(results[j]["locationtext"])
            d_stream = results[j]["station"]["downstream"]
            if d_stream == end_doc["station"]["stationid"]:
                depart = True
                break
            break
    if(depart == True):
        break

# Add end point
routes.append(end_doc["locationtext"])
```

```
(base) C:\gitRepo\CloudCluster\pdx-cs-cloud-cluster\queries>python3 query5.py
Johnson Cr NB
Foster NB
Powell to I-205 NB
Division NB
Glisan to I-205 NB
Columbia to I-205 NB
```

Appendix

Q6: Update: Change the milepost of the Foster NB station to 22.6.

```
results = col_detectors.find({"locationtext": "Foster NB"})
milepost = "18.1"

print("\n\n---- Display the data before updating milepost ----")
for r in results:
    print("milepost:", r["milepost"])
    col_detectors.update_one({"milepost":r.get("milepost")}, {"$set":{"milepost": milepost}})

results = col_detectors.find({"locationtext": "Foster NB"})

print("\n\n---- Display the data after updating milepost to " + milepost + " ----")
for r in results:
    print("milepost", r["milepost"])
```

Appendix

Data size:

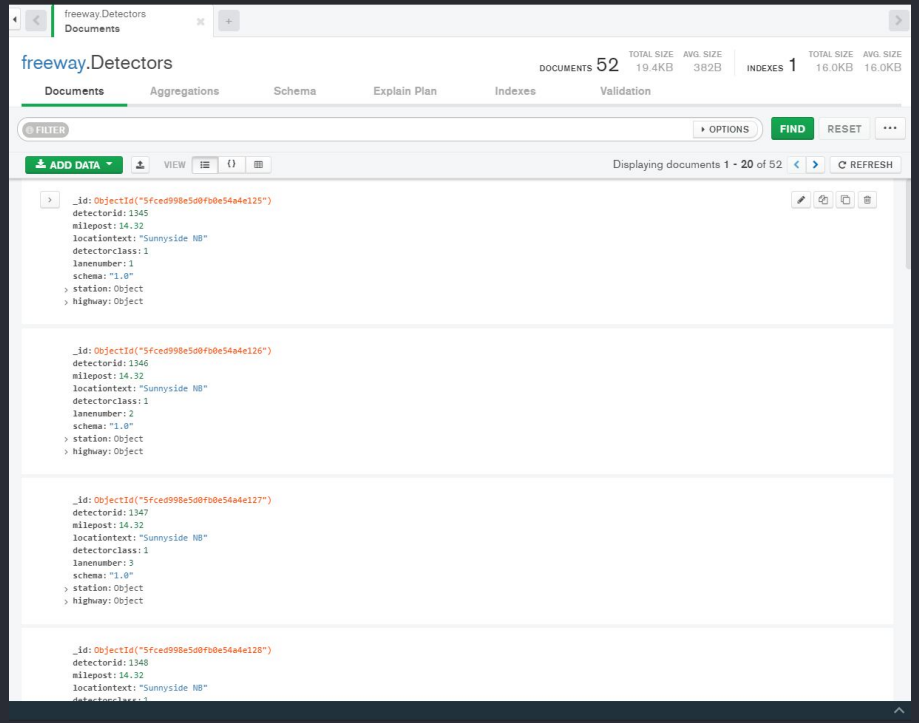

Appendix

Data size: